

BITS Pilani, Pilani Campus
2nd Sem. 2017-18
CS F211 Data Structures & Algorithms

Lab III - (15th Jan. to 20th Jan.)

Topics: Runtime Memory Layout and Heap Allocation, Linked Lists, Measuring Running Time and Space Usage

Programming: Con Linux

Memory Layout

- The typical runtime virtual memory (i.e. the collection of logical addresses) of a process (i.e. a program under execution) is divided into two parts: code area (which is typically read-only) and data area.
- The data-area is usually divided into three logical parts:
 - global / static area, meant for allocating global variables (referred to as static variables because they are allocated once at the start of program execution and stay allocated until execution ends)
 - stack, (a.k.a. call stack a.k.a. activation stack) where local variables of each function/procedure are allocated when that function is called and deallocated when it returns. (Local variables are referred to as automatic variables because they are allocated and deallocated automatically on call and return of functions/procedures).
 - heap, meant for dynamic allocation (i.e. allocation specified by programmer and executed at runtime) and deallocation. (In C, dynamic allocation and deallocation are explicitly specified by programmer e.g. by invoking malloc and free explicitly. In contrast, in Java, dynamic allocation happens when a new object is created, and deallocation is implicit – done by a garbage collector.).

Exercise 1: [Expected Time: 15 minutes.]

Write a program with multiple procedures (say, p, g, h, and d) invoked from main – both one after the other and from each other. Let each of these procedures contain a local variable pilani, goa, hyd, and dub respectively. Let bits be a global variable.

- (a) Print the addresses of the local variables pilani, goa, hyd, and dub under different call sequences. Use the unary prefix operator & to obtain the address and use "%u" format specifier (in printf) to print an address as unsigned int.
- (b) Modify the procedure p such that it calls itself with a single updated argument e.g. p(n) {...p(n+1)...}. Print the address of n each time p is invoked. Modify the test for termination of the recursion and repeat the test until you get a runtime error.
- (c) Summarize your understanding (for yourself) of the stack space allocated.

Exercise 2: [Expected Time: 25 minutes]

- (a) Define your own allocation and deallocation procedures myalloc and myfree respectively such that:

- myalloc invokes malloc to allocate the space as requested and returns the starting address of the allocated block; in addition, myalloc updates a global variable that keeps track of total space allocated in the heap so far.
 - myfree invokes free to free the space pointed to by the given argument, and in addition, updates the global variable that keeps track of total space allocated in the heap.
- (b) Write a loop that repeatedly allocates and frees a dynamic array of size M using your myalloc and myfree procedures. In each iteration:
- i. choose a random number M in the range 10,000 to 25,000.
 - ii. allocate an array A of M integers. Use sizeof to make it portable.
 - iii. print the addresses of the first and the last location of A i.e. A and &(A[M-1])
 - iv. free A
- The loop should terminate when allocation fails. Test the return value of malloc for failure.
- (c) Summarize your understanding (for yourself) of heap space used.

Linked Lists

A linked list is said to be:

- linear if traversing the list from the head ends in a “last node” which does not point anywhere i.e. there is a node whose next is set to NULL.
- cyclic if traversing the list from the head leads one to cycle the list i.e. one of the nodes points to another node in the list.
- circular if traversing the list from the head leads one to cycle the list through the head i.e. one of the nodes points to the head node. [Note that a circular list is a special case of a cyclic list.]

Exercise 3: [Expected Time: 45 minutes]

- (a) Write a procedure createList(N) that generates N random numbers and stores them in a linear linked list Ls and returns Ls. All the allocation in this procedure must use myalloc. Output the total heap space allocated to a text file. N must be large (say 1 million or more.)
- (b) Write a procedure createCycle(Ls) that tosses a coin – programmatically – to decide whether Ls must be linear or cyclic. If it must be cyclic, this procedure picks a random number, say r, and lets the last node in Ls point to (i.e. set its next to point to) the node containing r. If it must be linear this procedure returns Ls as is.
- (c) Write a main program that creates a new linked list Ls using the createList and createCycle procedures. Note that such a list Ls may or may not be cyclic. The main program should classify Ls as linear or cyclic by invoking a procedure testCyclic that is declared in a header file cycle.h. Test the program with a dummy implementation of testCyclic that returns FALSE always (or TRUE always).

Testing for Cycles in Linked Lists:

There are two ways to test for a cycle in a linked list:

1. Hare-and-Tortoise algorithm:

Maintain two pointers hare and tort such that hare jumps two nodes in the linked list when tort jumps one node. Repeat the jumps in lock-step until hare finds itself behind tort or in step with tort. Ensure that trivial/special cases are handled properly: empty list, singleton list, a list with two nodes, and a linear list (i.e. one that terminates.)

2. Link-Reversal algorithm

Reverse the links of a linked list as you traverse the linked list starting from the head. What will happen if there is a cycle in this case?

Exercise 4: [Expected Time: 60 minutes]

- a) Define two different implementations of `testCydic` in `cycle1.c` and `cycle2.c` using the two different techniques mentioned above.
- b) Compile these two files separately with the main program file (and other required files) to create two different executables.
- c) Run the executables multiple times for different values of N, the size of the linked list created.
 - o Measure the running time taken (programmatically) and output N and time taken along with the amount of heap space used.
 - o Measure the running time of individual procedures using the gnu profiler for each run.
 - o Plot the values and summarize your understanding of the running times of these procedures.

Exercise 5: [Expected Time: 30 minutes]

Write a procedure `makeCircularList(Ls)` which tests whether a given linked list is linear or cydic:

- If it is linear it creates a circular list by setting the last node of Ls point to the first node of Ls.
- If it is cydic it deletes all nodes from the head before the cycle i.e. after deletion, the result is a circular list and only nodes that were within the cycle remain.

`makeCircularList` should return the modified list. Use `myFree` to delete nodes. Output the total heap space used at different points.